

Evolutionary Computation

Introduction and overview

- EC models the process of evolution as found in nature (biological evolution).
- Charles Darwin vs Jean-Baptiste Lamarck
- Lamarckian view:
 - Inheritance
 - Adapt over lifetime
 - Use and disuse
- Darwinian view (EC is largely based on this view):
 - Natural Selection (The Origin of Species)
 - All species share common ancestors
 - Survival of the fittest (Fitness)
 - Also inheritance (because of reproduction)
 - Mutation

- Evolutionary Algorithms (EA) in AI is an umbrella term for all the algorithms that model biological evolution (evolve solutions) and is a subset of EC.
- EC Paradigms (EAs)
 - Genetic algorithms (GA) - models genetic evolution
 - Genetic programming (GP) - GA's cousin
 - Evolutionary programming (EP) - phenotypic evolution
 - Evolutionary strategies (ES) - evolution of evolution, with strategy parameters
 - Differential evolution (DE) - GA with a different reproduction system (crossover)
 - Cultural evolution - models the way in which culture (the general beliefs or preferences of a population) influences the evolution of a population (genetic and phenotypic)
 - Co-evolution - two or more "species" compete or work together, evolving towards a solution. Predator-prey relationship with competitive co-evolution vs cooperative coevolution.

Genetic Programming:

- John Holland, widely considered as the father of GA
- Holland was not the first to propose GA but he popularized the idea

- Genotypic evolution i.e. the genetic characteristics of individuals
- Key characteristics of GAs:
 - Selection
 - Reproduction (crossover)
 - Mutation
 - Fitness
- The survival of the fittest

Search/optimization method

- GAs are used to find solutions in some search space similar to the way a local search method works (How does local search work?).
- Formally: A population based, parallel search or optimization algorithm.
- We have group of proposed solutions which are then continually modified through the GA operators to obtain possibly better solutions.
- Optimization in terms of trying to “optimize” the *fitness* of proposed solutions
- Liken to function optimization (Example?)

Representation schemes

- How do we represent *solutions*?

- Solution = individual = chromosome
- Usually have more than one individual in a population
- The very first representation schemes where bitstrings (Holland)
 - 1010010111010
 - What can we do with bitstrings?
- Nowadays n-dimensional vectors with real valued elements are popular
- Can actually have n-dimensional vectors with mixed-valued elements
- Terminology:
 - *Chromosome*: A proposed solution produced and manipulated by the GA operators
 - *Gene*: An element of a chromosome. One of the variables in the problem domain.
 - *Allele*: A value for a gene at a particular point in time.
 - *Population*: A collection of chromosomes.
 - *Generation*: A population at a given point in time during the execution of the GA. Epochs vs Generations



Figure 1: An example chromosome

Fitness Function

- Once again: Survival of the fittest
- Fitness is a measure how good a solution is, i.e. the quality of a proposed solution.
- Almost like the heuristic function used in game trees but not quite the same thing.
- The better the solution represented by the chromosome, the fitter the chromosome is.
- Fitness is a function that we are trying to maximize or minimize.
- It is often (but not always) the distance in search space from some desired solution (e.g. Euclidean distance).
- A fitness function F takes a chromosome C and produces a fitness value for C .

$$- \textit{Fitness} = F(C)$$

- Often used to drive selection.
- Absolute vs Relative.
- Objective vs Subjective.
- Stopping criterion.

- Example?

Crossover

- Represents reproduction as found in nature.
- Recombines parent genes to create offspring.
- Applied at some probability of crossover.
- Offspring (usually) succeed parent chromosomes for the next generation.
- Types of crossover:
 - Asexual (one parent)
 - Sexual (two parents, most common)
 - Multi-recombination (three or more parents)
- How do we apply crossover?
 - Assume we selected 2 parents and generate 2 offspring
 - One-point crossover
 - Two-point crossover
 - Uniform (n-point) crossover
 - For bitstrings we generate a mask to do crossover.
- Example?

Mutation

- Small changes in genetic makeup to (hopefully) render an advantage to survive in the environment.
- Adaptability?
- As a rule, mutation is only applied to offspring.
- *Main purpose:* To introduce (and possibly maintain) genetic diversity in a population, i.e. explore more solutions.
- Prevents stagnation of a population.
- Exploration vs Exploitation.
- Usually applied at a low probability P , why?
- Better strategy for the probability of mutation?
- Which chromosomes do we want to mutate more?
- How is mutation applied?
 - Bitstrings.
 - Floating points.
 - Other?

Selection

- Very important, used for almost all the GA operators to decide on which chromosomes to operate.

- Very big friends with fitness function (fitness is more often than not used as a parameter in the selection process).
- Cat breeder?
- Used to:
 - Select parents to crossover.
 - Select chromosomes to mutate
 - Next population!!!
- Selective pressure:
 - The importance of fitness.
 - High vs Low selective pressure.
 - Exploration vs Exploitation.
 - Stagnation?
- Types of selection operators:
 - Random selection
 - Proportional selection
 - Roulette wheel selection (form of proportional selection)
 - * Danger?
 - * High selective pressure, why?
 - Tournament selection (chromosomes compete among one another)
 - * High or low selective pressure?
 - Rank-based selection
 - * Linear vs Nonlinear

- * Arranges chromosomes and then selects index i .
- * $i \sim U(0, U(0, n - 1))$, n = the number of chromosomes in the population.
- Elitism (best chromosomes survive to next generation).
- Hall of fame (over all generations, best solutions thus far)

Stopping conditions

- Can't execute indefinitely, we would like a solution at some stage.
- When can we stop?
 - Fitness, we found a solution!
 - Maximum generations.
 - Stagnation.
 - * No improvement in fitness (average of population or best) for a number of generations.
 - * The population becomes homogenous.

Pseudocode

1. Initialize the generation counter $G = 0$
2. Create and initialize a population P_0 , containing N chromosomes
3. While NOT(**stopping conditions**) do
 - a) Calculate the fitness of each chromosome C_i in P_G , $f(C_i)$ for $i = 1, \dots, N$
 - b) Create a new population O to contain the offspring and populate it by applying **crossover** to **selected** chromosomes from P_G
 - c) Apply **mutation** at some probability to the chromosomes in O , the offspring
 - d) Select the next population P_{G+1} from P_G and O , $P_{G+1} = \text{Select}(P_G, O)$
 - e) Increment the generation counter, $G = G + 1$
4. end while

- Chromosome representation?
- Initialization?
- Stopping conditions?
- Fitness function?
- Selection?
- Crossover?
- Mutation?

Example application

- Function optimization.
- Finding solutions to systems of equations.
- Neural Network training.
- Combinatorial problem solving in general.
- Data clustering.
- Solve N-queens problem.
- Routing, scheduling, planning (TSP which is NP-complete!).
- Picture evolution.
- Intelligent agents.
- Code breaking.
- MANY MORE!

Genetic Programming:

Intro

- Originally proposed by John R. Koza in the late eighties (1989).
- Used to evolve computer programs (e.g. S-expressions for LISP).
- A bit more computationally expensive than other EAs.
- Almost exactly the same thing as GAs, biggest difference is the representation scheme used.

Representation

- Trees.
- What can be represented as trees?
 - Mathematical expressions and functions.
 - Boolean expressions (AND, OR, NOT).
 - Programming languages.
 - Decision trees.
 - Game trees.
 - Graphics or pictures???
 - More?

- How do we read trees?
 - AND's (down the branch to a child node) and OR's (between sibling nodes)

Grammar

- Rules used to combine the “primitives” of a language into useful and meaningful constructs.
- Defines valid or legal ways to combine simple entities (words, punctuation, variables etc.) into more complex structures.
- Terminal set, function set and semantics.
- Example: Mathematical expressions
- Terminal set (corresponds to leaf nodes) = $\{1, \dots, 9, \text{all valid variable names e.g. } x, y \text{ and } z\}$.
- Function set (corresponds to internal or non-leaf nodes) = $\{+, -, *, /, \log, \ln, =, \text{etc}\}$.
- Then the grammar could be:
 - $\text{Expr} ::= \text{Binary} \mid \text{Unary} \mid \text{Terminal}$
 - $\text{Binary} ::= \text{Expr BINOP Expr}$
 - $\text{Unary} ::= \text{UNOP Expr}$
 - $\text{Terminal} ::= \{\text{variables and numbers}\}$
 - $\text{BINOP} ::= \{+, -, *, /\}$
 - $\text{UNOP} ::= \{\log, \sin, \cos, \tan \text{ etc.}\}$
- Note how naturally trees and grammars fit together.

Mutation and crossover

- Used for the same purposes in GP as in GAs
- Type checking becomes an issue.
- Mutation:
 - Growing
 - Shrinking (or truncation).
 - Node swapping.
 - Randomly altering node elements (beware non-terminal and terminal).
 - Replacements of nodes.
 - Which node?
 - Which operators?
- Crossover:
 - Any crossover technique as for GA.
 - Usually one-point crossover.
 - One offspring vs two offspring.

Fitness

- Size?
- Accuracy (depends on application and if we are evolving decision trees)
- Number of patterns covered...

Example applications

- Program evolution.
- Evolving decision trees (data mining).
- Planning.
- Find the function (function learning)
- www.genetic-programming.org for some more examples

Questions to think about

1. Discuss how you would use a GP in terms of:

- Tree representation
- Grammar
- Mutation
- Crossover
- Fitness

To solve the following problems:

- To find the expression (possibly polynomial) for a given graph.
 - To evolve a decision tree for some given dataset.
 - To alphabetically sort elements in some given list (hint: Binary Search Tree).
2. Propose a grammar (including terminal and non-terminal sets) to be used to evolve Boolean expressions.

3. Would it be a good idea to use GP to evolve a game tree?

EC for Data Mining:

GA and Rule Extraction:

- Representation:
- Michigan approach:
 - Individual = 1 rule.
 - Fixed length chromosomes.
 - Rule sets? Run multiple times...
- Pittsburg approach:
 - Set of rules for each individual.
 - Variable length individuals.
- Binary Strings:
- Length of string per feature = the number of values that the feature can take on.

Att1 = {blue,red,green}

Att2 = {yes,no}

Class = {male,female}

Att1 Att2 Class

101 10 0

Rule:


```
if (Att1 = blue or green)
    AND (Att2 = yes)
then
    class=male
```

- What about attributes with all 1/0?
- What about more than 2 classes?
- Mutation?
- Crossover?
- Fitness function?
- Other operations:
 - Generalizing crossover:
 - * OR between crossover points.
 - * Generalizes rules.
 - * Overfitting?
 - * Copy features...
 - Specializing crossover:
 - * AND between crossover points.
 - * Prevents underfitting
 - * Removes feature tests.
- What about different data types?

Clustering:

- Representation 1:

- Number of patterns = length of chromosome. Each gene corresponds to 1 pattern. Assign cluster labels to each gene.
- No crossover, does not make sense.
- Mutation?
- Fitness function?
 - * Find center of cluster (how)?
 - * Euclidean distance of all associated patterns from cluster center.
 - * Average over all clusters...
- Representation 2:
 - Chromosome = cluster centroids...
 - Know number of clusters in advance = fixed length individuals.
 - Evolve number of clusters = variable length individuals

Feature Selection:

- Each gene is one feature.
- Takes on the value 1 (use feature) or 0 (omit feature)
- Mutation?
- No crossover.
- Fitness? Determined by data mining algorithm.

GP:

- Evolve:
 - Decision Trees: Leaf nodes contain class labels.
 - Regression Trees: Leaf nodes are continuous values.
 - Model Trees: Leaf nodes contain models (usually expression trees/functions)
- What about using GP for hierarchical clustering?

References:

- AP Engelbrecht, Computational Intelligence: An Introduction, Second Edition, Wiley & Sons, 2007.
- I De Falc, A Iazzeta, E Tarantino, AD Cioppa, An Evolutionary System for Automatic Explicit Rule Extraction, CEC2000.
- CB Congdon, Classification of Epidemiological Data: A Comparison of Genetic Algorithm and Decision Tree Approaches, CEC2000.
- JJ Liu, JT-Y Kwok, An Extended Genetic Rule Induction Algorithm, CEC2000.
- WM Spears, DF Gordon, Adaptive Strategy Selection for Concept Learning.
- KA De Jong, WM Spears, DF Gordon, Using Genetic Algorithms for Concept Learning.